

I Missed That

Transcription tells you what was said.

We tell you what you missed.

A live decision & commitment ledger for deaf and hard-of-hearing users at the table.

The pain isn't the words. It's what's around them.

"I work in marketing and we have these weekly team meetings where everyone just talks over each other... I've started just nodding along even when I have no idea what was decided. My boss mentioned something about a new client project last Thursday and I smiled and agreed but I genuinely don't know what I signed up for."

— r/deaf

"Not hearing the whole joke or the punchline in the moment at a table full of people so everyone is laughing but you, so you have to either ignore it, fake laugh, or ask them to repeat it while everyone is silent. Totally awkward."

— r/hardofhearing

Why "just add captions" isn't enough

- **51.9%** of DHH caption users are frustrated with live captioning quality
- ASR lag of **1–2s** is rated *worse* than the errors themselves
- DHH caption comprehension drops above **170 wpm** — live group speech runs **160–220 wpm**
- Only **4%** of non-speech info (laughter, tone, applause) ever reaches a caption
- **96%** of deaf children are born to hearing parents — Dinner Table Syndrome is the default, not the exception

Sources: [pitch/citations.md](#)

What we built


Three static cards, not a transcript:

Now	Open on the table	For you
who's talking + last line	decisions forming, objections, open questions	anything just addressed to you

Plus **Catch me up** (2–3 bullets since you looked away) and **laughter recovery** ("they laughed at Sam's Nokia line").

State, not words. We track what changed, who owns what, and whether you were just asked something — not a wall of scrolling text.

How it works



Diarized speech and non-speech events feed one shared timeline; Claude turns it into state, never prose.

What's next — and the honest limits

Next: persistent speaker memory across sessions, on-device (no-cloud) mode, wearable haptic nudge, multi-language diarization.

Limits, honestly:

- Diarization still flaps on heavy crosstalk — we bias toward "someone" over a wrong name
- No Deepgram key → falls back to browser ASR, captions only, no speaker colors
- Nothing is stored after the session (privacy by design) — also means no long-term history yet
- Built and tested in 12 hours — this is a proof of the *idea*, not a hardened product